

PRODUCT OVERVIEW

Evolve AI Enterprise Delivery Studio

Take a client from the first conversation to a deployed, auditable AI solution, in one desktop app that can run entirely on your own hardware.

- Windows desktop app
- VS Code extension
- Local or cloud AI
- Air-gap capable

MAP THE DATA

Table

Select any table below or tap nodes on canvas to inspect schema and records.

Quick filter list... A-Z

addresses

173,892 rows • 7 cols • 4 FKs

categories

140,107 rows • 4 cols • 1 FKs

customer_reviews

88,325 rows • 6 cols • 6 FKs

customers

3 FKs

inventory_transactions

order_items

2 more columns

orders

4 more columns

payments

2 more columns

products

3 more columns

shipments

2 more columns

2D/3D

2 - THE 8 SPECIALIZED AI MODELS (AUDITED AGAINST LITERATURE)

LLM

Autoregressive next-token text generation & instructions

Job: chat

Inspect --

LCM

Sentence concept space modeling vs. few-step diffusion

Job: reasoning

Active

LAM

Marketing Coinsage -- Tool Agent

Autonomous action execution & MCP function calling

Job: code-agentic

Inspect --

VLM

Multimodal image, document & diagram understanding

Job: vision

Inspect --

SLM

Air-gapped edge & laptop deployable (Size Attribute)

Job: chat

Inspect --

MLM

Pre-Training Objective -- Encoder

Matched

Job: embedding

Matched

SAM

Interactive promptable zero-shot vision segmentation

Job: vision

Inspect --

LCM (Large Concept Model / Latent Consistency)

2 Fused Papers (Audited)

Job: reasoning

ollama pull deepseek-r1:8b

Git & Remote Repository Hub

Bitbucket Cloud - 7 Changes

Open Local Repository

SHIP IT

1-Click Commit & Push

Create PR

Setup wizard

Terminal

ACTIVE BRANCH

main

REMOTE PROVIDER

Bitbucket Cloud (retail-pilot.git)

WORKING TREE

7 Modified Files

GIT AUTHOR

EvolveMinds (support@evolveminds.com.au)

Working Tree Changes

7 Files

Stage All

Reset

.evolve/

Added / Untracked

.github/

Added / Untracked

docker-compose.yml

Added / Untracked

k8s/

Added / Untracked

Commit & Upstream Push

AI Generate Message

feat(module): describe your changes cleanly...

THE PROBLEM

Delivering AI to a client is still stitched together by hand

The work between "we have a use case" and "it's running in production" is spread across a dozen tools, and most of the judgement calls live in people's heads.

01

Context lost at every hand-off

Schema notes, notebooks, diagram tools, the IDE and three cloud consoles. Each move between them drops a decision on the floor.

02

Data that can't leave the building

Banks, government and healthcare can't send schemas or source code to a public model. Most AI tooling assumes they can.

03

AI where a rule would do

Teams reach for a language model by default, even where a deterministic rule would be faster, cheaper and provably correct.

EVOLVE AI

Puts the whole engagement in **one workspace**, runs on **local models** when the data can't leave, and tells you **when not to use AI at all**.

One workspace that follows the whole engagement

PHASE 1
Discover & Frame
Scope, ask, success
criteria


PHASE 2
Engineering Core
Client data, APIs, schemas,
marts

PHASE 3
AI Solutioning
Rule vs model, RAG, MCP

PHASE 4
Reliability & Evals
Test it before the client
does

PHASE 5
Deploy & Influence
Ship it and make the case


PHASE 6
DevOps & Git
Branch, commit, PR,
pipeline



See the client's data model in minutes, not workshops

- **Connect a live database or warehouse** (PostgreSQL, Snowflake, BigQuery and more) and pull the real schema.
- **Switch between a technical ERD and a 3D topology** of the same tables and relationships.
- **Go straight from diagram to build:** staging models, join paths and dimensional marts.
- **Export a POC approval pack** the client can sign off.

Credentials default to session-only, in memory.
Nothing is written to settings files.



It asks "do we need AI?" before anyone builds any

- **Rule-vs-model decision gate.** Describe the task: its arithmetic, inputs, latency budget and tolerance for error.
- **A five-level capability ladder** recommends the simplest architecture that does the job, from compiled rules up to agentic systems.
- **Every recommendation comes with its reasoning**, a flow diagram, a trade-off matrix and a code contract.
- **Export architecture decision records** straight into the client's documentation.

For a ledger reconciliation it recommends a rule engine, not a language model. That's the answer a trusted advisor gives.

Phase 3 · AI Solutioning ("FDE 1–5 Ladder, Rule vs Model, RAG & MCP")

Air-Gapped Ready

Architect solutions using the canonical FDE 1–5 capability ladder. Enforce deterministic rule-versus-model decision gates and scaffold air-gapped hybrid policy RAG and MCP tools.

SOLUTION CONTRACT: FinOps & Banking → Pure Rule Engine / Compiled SQL → Level 1 → MLM (nomic-embed) + SLM (Qwen 2.5 7B)

<5ms P99 • 0.0% (Zero Drift Guarantee)

3A Problem & Decision ...

3B 1–5 Capability Ladder...

3C Component Architect...

3D Solution Contract & S...

← Back

Next →

Section 3A: Problem Definition & Rule vs. Model Decision Gate ("Do We Need AI?")

DETERMINISTIC & HYBRID READY

Export ADR Docs

Ver: v1.0.0 (Evaluator Baseline) ▼

Snapshot

Edit Gate

Evaluate Gate

Save to Project

Decision Gate AI Co-Architect: (Natural language transformations & hybrid recommendations)

Suggest Hybrid

e.g. 'Convert to a Level 1 + Level 3 hybrid with strict SQL tolerance checks before RAG retrieval', 'Optimize for <15ms latency', 'Add strict SOX guardrails'...

Apply

Quick Prompts: Hybrid Level 1+3 Guardrailed Tool Agent (1+4) Swarm with HITL (1+5) Low Latency (<10ms) Challenge for SOX Audit

Enterprise Task Archetype Preset: Financial Tolerance & Ledger Balance (Level 1: Pure Rule Engine & SQL) ▼

Architecture Paradigm Override: Auto (Evaluator Recommended) ▼

Task Description	Arithmetic?	Input Modality	Latency Budget	Hallucination SLA	HITL Escalation
Transaction tolerance reconciliation and invoice	Strict Math (Zero I ▼)	Structured SQL / I ▼	50ms	0.0% (Zero Hallucin ▼)	Discrepancy Trigger ▼

Recommended Architecture: Pure Rule Engine / Compiled SQL (LEVEL 1)

LEVEL 1

SLA CONFIDENCE: 99.9%

Strict arithmetic calculations and monetary ledgers must never use non-deterministic probabilistic LLMs. Executed via deterministic compiled TypeScript/SQL boundary rules (<5ms latency, 0% hallucination SLA).

P99 LATENCY SLA<5ms P99

UNIT ECONOMICS\$0.0000 / op (Zero LLM Tokens)

HALLUCINATION SLA0.0% (Zero Drift Guarantee)

HITL TRIGGERHard threshold ceiling

Evolve AI · Product overview

5

The right retrieval pattern, then the right kind of model

- **Start from the client's failure.** Pick the symptom ("exact part numbers get missed") and it matches one of eight canonical RAG patterns, each with a diagram, trade-offs and a delivery playbook.
- **A model guide, fact-checked against the research.** The popular "8 types of AI model" list is checked against the published papers. It shows which types are real categories, which are marketing names, and which are really attributes such as size or Mixture-of-Experts.
- **Every model type links to the job it does** (chat, embeddings, vision, tool use), its architecture diagram and its primary paper.
- **Then it checks your machine:** which installed models can actually do that job.

Recommendation: Hybrid RAG selected. Fuses sparse BM25 lexical keyword search with dense HNSW vector embeddings to prevent vocabulary drift on product codes.

Select Architecture Paradigm:

Active: 06 Hybrid RAG (Production Gold Standard)

01

Naive RAG

Retrieve once, then generate.

<60ms · Low Cost

02

Multimodal RAG

Cross text, diagrams, audio, video.

Vision & ColPali Embeds

03

HyDE

Hypothetical doc guides retrieval.

Zero Vocabulary Mismatch

04

Corrective RAG

Evaluate evidence before using it.

Evaluator & Search Fallback

05

Graph RAG

Connected facts & graph summaries.

Entity Knowledge Graphs

06

Hybrid RAG

BM25 exact keywords + Vectors.

Enterprise Production Gold

07

Adaptive RAG

Dynamic complexity routing.

Optimized Cost & Latency

08

Agentic RAG

Plan, retrieve, inspect & repeat.

Cross-Tool & SQL Reasoning

2 · THE 8 SPECIALIZED AI MODELS (AUDITED AGAINST LITERATURE)

Click any model card to inspect architecture flow and test compatibility

LLM

Autoregressive next-token text generation & instructions

Job: chat

Inspect →

✓ Sound Category

LCM

Sentence concept space modeling vs. few-step diffusion

Job: reasoning

Inspect →

⚠ 2 Fused Papers (Audited)

• Active

LAM

Marketing Coinage → Tool Agent

Autonomous action execution & MCP function calling

Job: code-agentic

Inspect →

MoE

Internal sparse layer routing (Architecture Attribute)

Job: chat

Inspect →

✗ Category Error

VLM

Multimodal image, document & diagram understanding

Job: vision

Inspect →

✓ Sound Category

SLM

Air-gapped edge & laptop deployable (Size Attribute)

Job: chat

Inspect →

⚠ Deployability Filter, Not a Job

MLM

Pre-Training Objective → Encoder

Bidirectional attention embeddings & fast classification

Job: embedding

Inspect →

⚡ Matched

⚡ Matched

SAM

Interactive promptable zero-shot vision segmentation

Job: vision

Inspect →

✗ Product Name

LCM (Large Concept Model / Latent Consistency)

⚠ 2 Fused Papers (Audited)



Job: reasoning

ollama pull deepseek-r1:8b

"Sentence concept space modeling vs. few-step diffusion"

From raw data to a client-ready report, locally

- **Bring any source:** CSV, Excel, JSON and Parquet files, or BigQuery, Databricks, Cosmos DB, DynamoDB, cloud storage and SQL databases.
- **Choose the deliverable:** an HTML report, a reproducible notebook, a data profile, or interactive 3D analytics.
- **Explore in 3D** with scatter, PCA, surface and k-means views, and outliers flagged automatically.
- **Refine in plain language.** "Move the churn chart up" edits the report; every change can be undone.



Move legacy code to a modern stack, and keep an audit trail

- **26 languages**, including COBOL, VBA, SAS and Perl on the legacy side. Convert a selection, a file or a whole folder.
- **Choose the style:** idiomatic, line-for-line literal, or modernised, with a dependency policy to match.
- **Checks its own work** with the target language's own parser where it's installed, and feeds failures back for repair.
- **Every conversion comes with a report** of what mapped exactly, what was approximated and what still needs a person.

Nothing is written to disk until you accept the side-by-side review.

LEGACY & ENTERPRISE

Perl .pl
VBA .vba
COBOL .cbl
SAS .sas

Optional — target framework or runtime, e.g. "FastAPI", ".NET 8", "Spring Boot", "React"

3 - TRANSLATION FIDELITY

Idiomatic (Recommended)
Target-language idioms, standard library conventions, clean modern code

Literal
Line-by-line mechanical translation preserving original architecture and variable names

Optimized
Refactored for performance, modern concurrency, and vectorized operations where available

4 - DEPENDENCIES & TESTS

Popular Ecosystem Packages
Uses standard idiomatic packages (e.g. axios/requests, serde, pydantic)

Standard Library Only
Zero external packages — uses strictly built-in runtime modules

Preserve Equivalent Libraries
Maps 1:1 to direct package equivalents (e.g. lodash → ramda)

5 - AI ENGINE & MODEL FIT

local / provider - qwen2.5-coder:7b
RECOMMENDED LOCAL CODING (AIR-GAPPED)
Switch Model / Provider

Specialized coding model with 92%+ pass rate. Fast, offline, and secure AST translation to TYPESCRIPT with zero data egress.

Comfortable Local Fit: Air-gapped offline conversion without sending code to cloud endpoints.

6 - OUTPUT DESTINATION & STORAGE LOCATION

DESTINATION FOLDER PATH:
src/converted
Browse Folder...

TARGET FILE NAME:
converted_output.ts

Target Destination: src/converted/converted_output.ts
Auto-creates subdirectories if missing

☐ Also generate unit tests
☒ Carry comments across
☒ Emit dependency manifest

Anything else the conversion must respect — e.g. "keep the CLI flags identical", "must run on Java 11", "no external HTTP client"

Convert Code →

From architecture to deployable infrastructure in one click

- **Pick the target:** Google Cloud, AWS, Azure, or on-prem and air-gapped Docker.
- **Size it from a preset,** from a scale-to-zero pilot sandbox up to GPU inference, then fine-tune CPU, memory, network and autoscaling.
- **One click scaffolds everything:** Terraform, Kubernetes, Docker Compose, deploy scripts for bash and PowerShell, and a CI/CD workflow.
- **Lists the exact secrets** the pipeline needs, and targets GitHub Actions, GitLab CI, Bitbucket or Azure DevOps.

Google Cloud (GCP/Firebase) Amazon Web Services (AWS) Microsoft Azure On-Prem / Air-Gapped

Discover & Auto-Fill (GCP API)

Environment Sizing & Architecture Presets
1-click configuration for development sandboxes, high-availability clusters, GPU inference, or air-gapped on-prem.

Preset: **Pilot Sandbox (Scale-to-Zero, 1 CPU, 1Gi, Direct HTTPS)**

Profile: Serverless Managed Container · Autoscale 0-10 instances · Single Region · Public HTTPS

Google Cloud Deployment Parameter Matrix

CPU Allocation (Cloud Run / GKE)	Memory Allocation (RAM)	Hardware
1	1Gi	None (\$)
Client VPC / Network ID	Subnet ID / Subnetwork	Security Group
e.g. default or custom-vpc	e.g. pilot-subnet or subnet-0a1b2c	e.g. allow-https
Network Ingress Security	Autoscaling Bounds (Min - Max)	Secrets & Key Management
Public HTTPS (Direct)	0 10	Google Cloud KMS

1-Click Infrastructure Scaffolding (Terraform / K8s / Docker / Cross-Platform Runners)

Target Cloud Project ID / Account: acme-pilot-2026 Deployment Region / Location: australia-southeast1

Scaffold All Deployment & IaC Assets Generate Terraform (main.tf) Generate Kubernetes (k8s.yaml)

Interactive Multi-File IaC Result Drawer
Files ready for workspace: terraform/, k8s/, docker-compose.yaml, scripts/, .github/

terraform/main.tf k8s/deployment.yaml docker-compose.yaml scripts/deploy.sh

scripts/prepare-deployment.js .github/workflows/deploy.yml

Workspace Target: terraform/main.tf

```
# -----
# Scaffolding: Evolve AI (Forward Deployed Engineer Suite)
# -----

terraform {
  required_version = ">= 1.5.0"
  required_providers {
    google = {
      source  = "hashicorp/google"
      version = ">= 5.0"
    }
  }
}
```

Interactive Multi-File IaC Result Drawer
Files ready for workspace: terraform/, k8s/, docker-compose.yaml, scripts/, .github/

terraform/main.tf k8s/deployment.yaml docker-compose.yaml scripts/deploy.sh

scripts/prepare-deployment.js .github/workflows/deploy.yml

Workspace Target: terraform/main.tf

```
# -----
# Scaffolding: Evolve AI (Forward Deployed Engineer Suite)
# -----

terraform {
  required_version = ">= 1.5.0"
  required_providers {
    google = {
      source  = "hashicorp/google"
      version = ">= 5.0"
    }
  }
}
```

Cloud Provider Secret Requirements Matrix
Exact environment secrets required by CI/CD runners to deploy to Google Cloud (GCP/Firebase).

Secret / Env Name	Target CI/CD Secret Scope	Purpose & Expected Value
GCP_SA_KEY	Repository Secret	Base64 Service Account JSON with Cloud Run Admin & Storage roles
GCP_PROJECT_ID	Repository Variable	Target Google Cloud Project ID (e.g. acme-pilot-2026)
GCP_REGION	Repository Variable	Compute Region (e.g. australia-southeast1, us-central1)

Automated CI/CD Infra Pipeline & Git Push Hub
Generate ready-to-run CI/CD pipelines (GitHub Actions, GitLab CI, Bitbucket, Azure DevOps) with simulated verification.

CI/CD Platform: GitHub Actions (.github/workflows/deploy.yml) Deployment Target Branch: main Deployment Strategy: Pilot / Full

Scaffold CI/CD Pipeline File Simulate CI/CD Pipeline Run Run Deploy Script (Terminal) 1-Click Deploy


Simulated CI/CD Pipeline Visualizer

1 2 3 4

Git set up and shipping without leaving the tool

- **One-click setup wizard** for Bitbucket, GitHub and GitLab. It detects git, your identity and the remote, then initialises the repo and tests the connection.
- **See every change live:** branch, remote and working tree, updated in real time.
- **AI drafts the commit message**, then commit and push in one click.
- **Branch, pull, stash and open a pull request** from the same screen, with a git command bar for everything else.

The client's delivery work lands in their repository with a clean history, not on a consultant's laptop.



Runs on the hardware you already have

- **Detects the machine's RAM, GPU and CPU** and shows which local models will actually fit, before anything is downloaded.
- **One-click install** of a recommended local model.
- **Switch provider per job:** a bigger model for a code port without changing the default for everything else.

LOCAL & OFFLINE

Ollama, Gemma 4, GLM, Colibri, and a fully offline mode

CLOUD

Anthropic, OpenAI-compatible, Google Gemini, Z.ai, Hugging Face

Local AI Hardware Sizer & Model Engine Prober

+ Add Custom Model

Re-Scan System Hardware

Inspects host RAM, GPU VRAM, CPU cores, and probes active Ollama (11434), LM Studio (1234), and vLLM (8000) servers.

SYSTEM MEMORY (RAM)

16 GB

Host physical memory

DEDICATED GPU / VRAM

NVIDIA GeForce RTX 3060 Ti (8GB VRAM)

VRAM Acceleration

CPU COMPUTE

16 Cores

AMD Ryzen 7 5800 8-Core P...

LOCAL AI SERVER

✓ Ollama

Port 11434 / 1234 / 8000

Recommended Offline Local Models for This Machine

Recommendation: gemma4:e4b (Best balance of quality and speed for your system)

Calibri 744B Feasibility: Not enough disk space — GLM-5.2 INT4 needs about 372GB

Local On-Premise AI Models & Hardware Compatibility Matrix

Search local models...

All Specs (23)

Compatible Only

Installed

Evaluated against your active host GPU VRAM and System RAM. Compatible models can be installed with 1 click.

All Providers

DeepSeek

Google Gemma

GLM / Z.ai

Alibaba Qwen

Meta LLaMA

Mistral & MS

Servers (LM Studio / vLLM)

Qwen 2.5 Coder 7B

Alibaba Qwen

Optimal GPU Fit (5GB / 8GB VRAM ~ ~60+ tok/s)

Premier 7B coding model with 92%+ pass rate. Highly optimized for SQL, Python, TypeScript migrations.

Min RAM: 8 GB Target VRAM: 5 GB Disk Footprint: 4.7 GB Context Limit: 32k

ACTIVE ENGINE

Qwen 2.5 Coder 14B

Alibaba Qwen

Compatible (CPU/GPU Hybrid ~ ~25-35 tok/s)

High-capability coding model matching GPT-4 on code refactoring and data engineering tasks.

Min RAM: 16 GB Target VRAM: 9.5 GB Disk Footprint: 9 GB Context Limit: 32k

Installed Activate

Qwen 2.5 Coder 32B

Alibaba Qwen

Exceeds Machine Specs (Requires 32GB RAM / 20GB VRAM)

Frontier-grade open coding model for full-repository modernization and zero-shot architecture design.

Min RAM: 32 GB Target VRAM: 20 GB Disk Footprint: 19.5 GB Context Limit: 32k

Installed Activate

Qwen 2.5 72B Instruct

Alibaba Qwen

Exceeds Machine Specs (Requires 64GB RAM / 44GB VRAM)

Heavyweight 72B open foundation model matching frontier closed models in polyglot reasoning.

Min RAM: 64 GB Target VRAM: 44 GB Disk Footprint: 41 GB Context Limit: 128k

Installed Activate

Gemma 4 e4b (Multimodal)

Google Gemma

Optimal GPU Fit (5GB / 8GB VRAM ~ ~60+ tok/s)

Google's newest open multimodal architecture for edge devices with exceptional efficiency.

Min RAM: 8 GB Target VRAM: 4.5 GB Disk Footprint: 4.2 GB Context Limit: 32k

1-Click Auto Install

Evolve AI · Product overview

11

AT A GLANCE

A complete delivery platform, shipping today

6

Delivery phases

Discovery through to DevOps, in one workspace

26

Programming languages

For review-first code conversion

24

Stack plugins

Databricks, dbt, Airflow, Terraform, Kubernetes and more

10

AI provider modes

Local, cloud or fully offline

8

RAG architecture patterns

Each with a diagram, trade-offs and a playbook

2

Surfaces

Windows desktop app and VS Code extension

Built for teams that answer to auditors

Air-gap capable

Run on local models with no network at all. Client data and code never have to leave the machine.

No invented numbers

Every figure in a client document is labelled as measured, user-supplied, demo or unmeasured. The product won't fill a gap with a plausible guess.

Review before anything changes

Generated code, conversions and edits are previewed before they land, and can be undone.

Secrets stay in the OS vault

API keys and access tokens live in the operating system's secure storage, never in settings files.

OUR RULE A visible gap is better than a plausible guess. **If nobody measured it, the product says so.**

WHO IT'S FOR

AI delivery that holds up in front of the client

Solutions & forward-deployed teams

One place to run an engagement from discovery to deployment.

Data & AI consultancies

Repeatable, well-documented delivery across many clients and stacks.

Regulated enterprises

AI assistance that stays inside the perimeter and leaves an audit trail.

Let's show you a live walkthrough · support@evolveminds.com.au